



Datasets and slides:

<https://github.com/doktor-nick/adv-ML-image-segmentation-with-UNET>

# Advanced Machine Learning Image Segmentation with UNET



**Dr Nick Hamilton**  
**Institute Bio-Mathematician**  
Institute for Molecular Bioscience  
Research Computing Centre  
Queensland Cyber Infrastructure Foundation  
[n.hamilton@imb.uq.edu.au](mailto:n.hamilton@imb.uq.edu.au)



[www.imb.uq.edu.au](http://www.imb.uq.edu.au)

Power corrupts.

PowerPoint corrupts absolutely.

# Outline

- Brief machine vision applications overview
- Semantic image segmentation
- The UNET architecture for image segmentation
- Setting up Google Colab
- Exercises
  - Introduction for UNET
  - More advanced UNET use


# Applications of Computer Vision

- Classification



Rainbow bee-eater

- Object Detection



Bee-eater, bee

- Semantic Segmentation




Bee-eater, bee, stick,  
background

- Instance Segmentation



Bee-eater 1 and 2

# Image Classification



=

```
( 0.000772532, 0.00154506, 0.00343348, 0.0139056, 0.0308155,  
 0.0440343, 0.0337339, 0.0433476, 0.828412, 0.000235031,  
 0.000352547, 0.00205653, 0.00593454, 0.0223867, 0.0351372,  
 0.022093, 0.0312004, 0.880604, 0.000420663, 0.00117786,  
 0.00496382, 0.0165741, 0.058388, 0.107017, 0.0904425,  
 0.110803, 0.61024)
```




Cell  
Cytoskeleton


Generate  
statistics

Machine  
learning  
classifier

# Threshold Adjacency Statistics



ER




Microtubules

Let  $\mu$  be the **mean intensity** in the image  
(Band pass) **threshold intensity** in range  $(\mu, \mu+30)$



# • Threshold Adjacency Statistics (TAS)



- Count number “on” pixels with

- 0 on neighbours
- 1 on neighbour
- ...
- 8 on neighbours



- Nine image statistics

# Image Segmentation



# Semantic Image Segmentation

## Definition

A segmentation is a partition of an image into regions that label the class of each region. e.g. foreground/background

## Approach

Instead of a whole image we want to label/classify individual pixels




Ureteric tree  
Cap mesenchyme  
Stroma  
Background

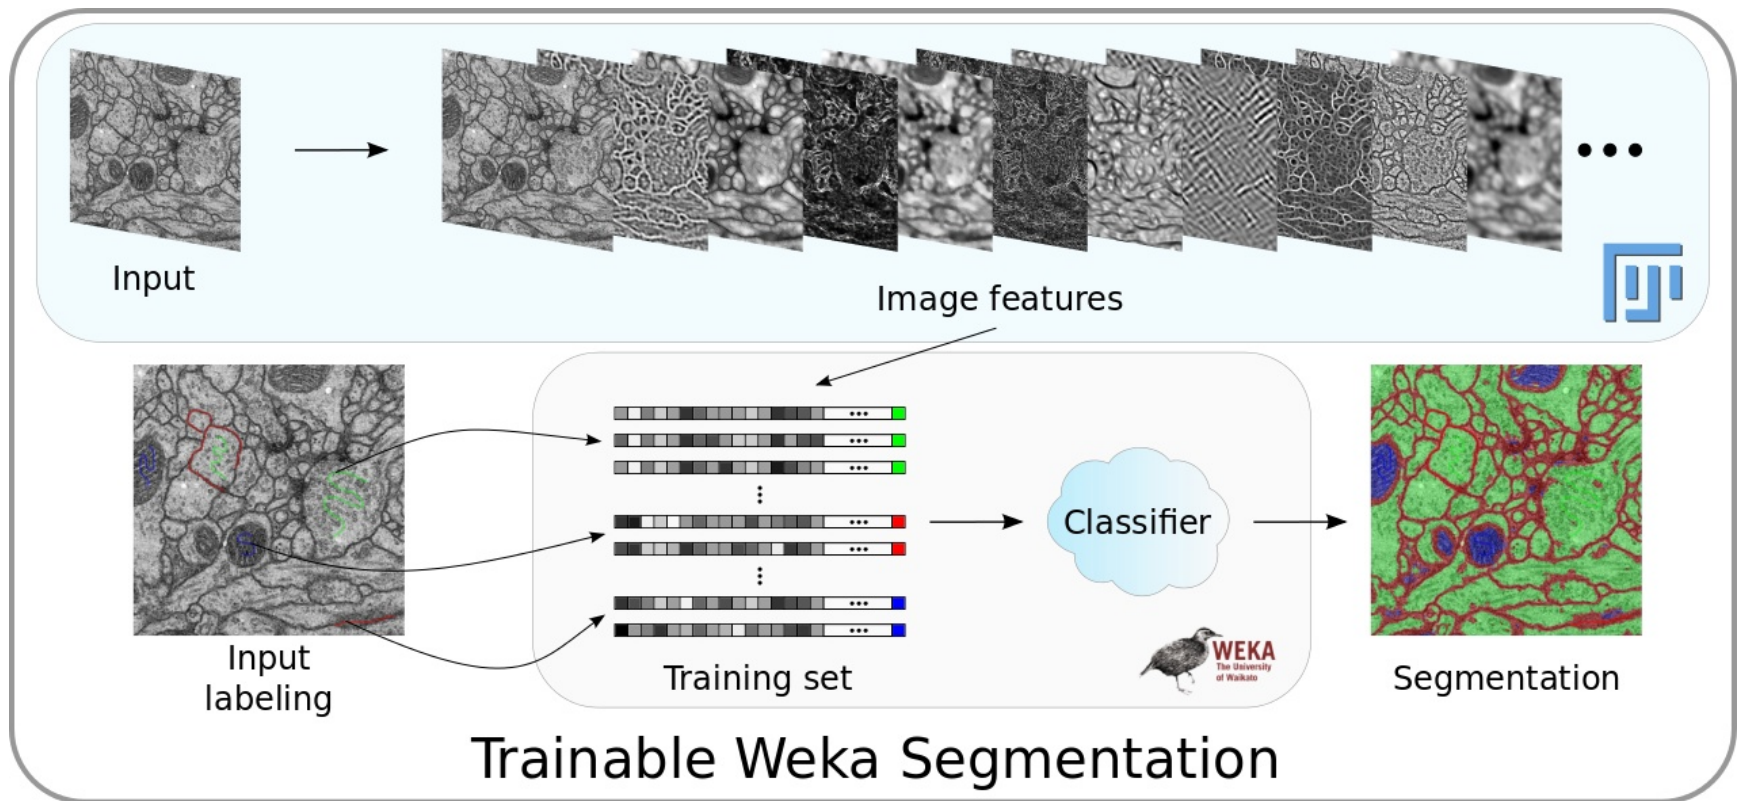


# Basic Image Segmentation

- Generate **statistics** for **each pixel** in the image
- E.g. intensity, local median intensity, gradient, ...
- Classify that pixel using machine learning



# The ImageJ/Fiji Trainable Weka Pipeline



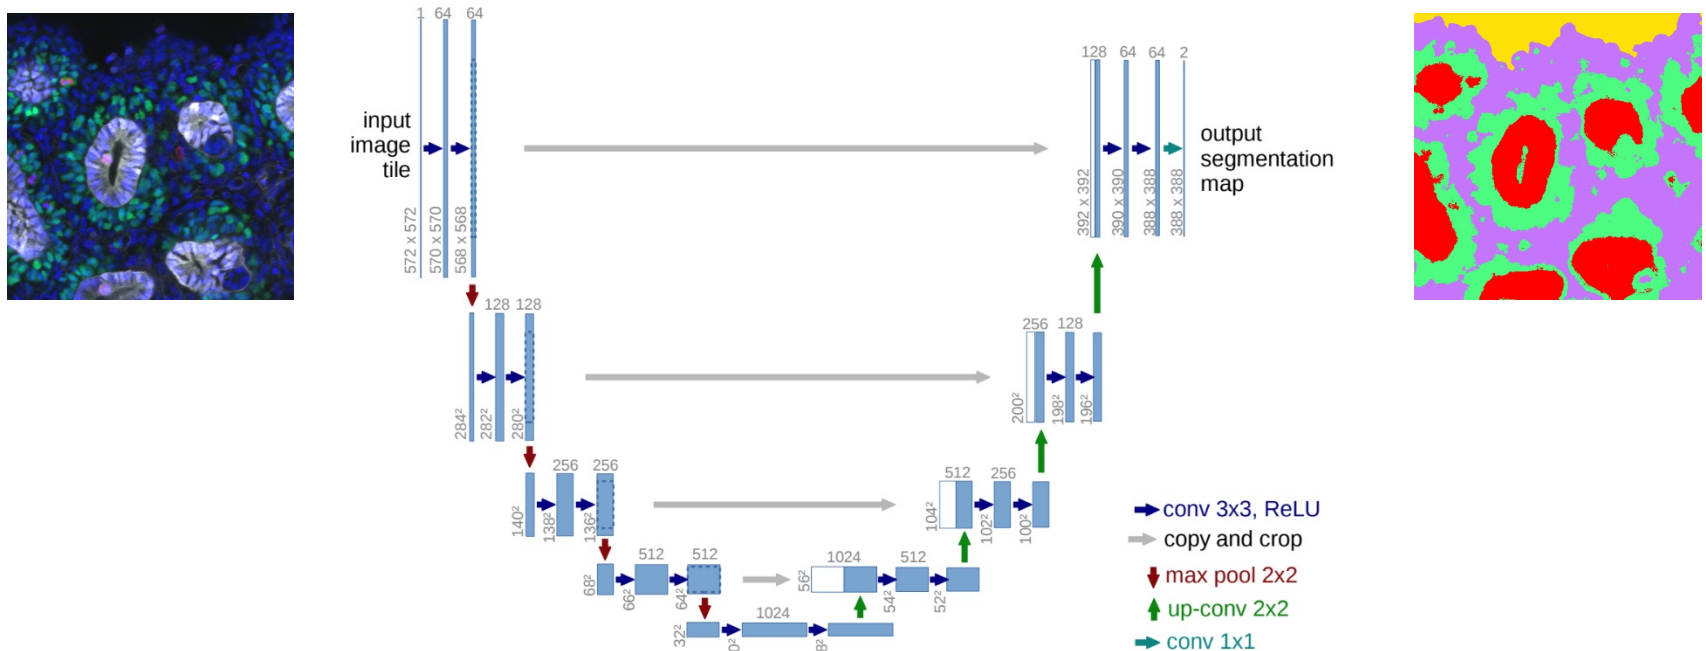
See “Introduction to Machine Learning for Imaging” workshop

Image: [https://imagej.net/Trainable\\_Weka\\_Segmentation](https://imagej.net/Trainable_Weka_Segmentation)

# The UNET Architecture for Image Segmentation

# Convolutional Deep Learning with U-net

- **U-net** was a new type of neural network architecture introduced in 2015
- It down-scales then up-scales an image to create a segmentation




- U-Net and its variants are the state of the art in segmentation
- U-net paper: <https://arxiv.org/pdf/1505.04597.pdf>

# Principle UNET Components

- 3x3 **Convolution** followed by ReLU (rectified linear activation function)
- 2 x 2 **Maxpooling** with stride 2 (halves image size)
- Double number of feature channels (convolutions) 64, 128, 256, 512, at each down sampling to compensate
- **Dropout layers**
- **Skip layers**
- **Transpose convolutions** / up convolutions

# Convolutions



$$Output = \sum Kernel_{i,j} * Input_{i,j}$$

- The kernel represents the pattern to be detected
- The more the input matches the kernel, the more positive the output response would be
- Convolutions are pattern detectors
- The kernel weights are learned during training  
i.e. it learns what patterns to detect and respond to



# A 3x3 Convolution Max

- **Convolution** of an image by a 3x3 matrix replaces each pixel in an image by a weighted sum of those pixels adjacent to it, i.e.

$$I'(x,y) = a I(x-1,y+1) + b I(x,y+1) + c I(x+1,y+1) \\ + d I(x-1,y) + e I(x,y) + f I(x+1,y) \\ + g I(x-1,y-1) + h I(x,y-1) + i I(x+1,y-1)$$

Local Image Region

$I(x-1,y+1)$	$I(x,y+1)$	$I(x+1,y+1)$
$I(x-1,y)$	$I(x,y)$	$I(x+1,y)$
$I(x-1,y-1)$	$I(x,y-1)$	$I(x+1,y-1)$

Convolution matrix

a	b	c
d	e	f
g	h	i

# Convolution Example

Convolution

0	0	0
0	1	1
0	1	1

Image

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	1	1	0	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

Image x Convolution

Dark	Light	Dark	Black	Black
Light	White	Light	Black	Black
Dark	Light	Dark	Black	Black
Black	Black	Black	Black	Black
Black	Black	Black	Black	Black

# Max Pooling with step size 2

10	6	6	0
5	-3	14	-2
0	2	-7	-3
0	0	-5	-5

Input 4x4 image



10	14
2	-3


Output 2x2 image

- Slide a window across the input and pick a value at every window position.
- Max pooling – take the max value.
- Average pooling – take the average value.
- Pooling layers are information filters.
- Images are reduced in scale, typically by  $\frac{1}{2} \times \frac{1}{2}$

# Spatial Dropout Layers

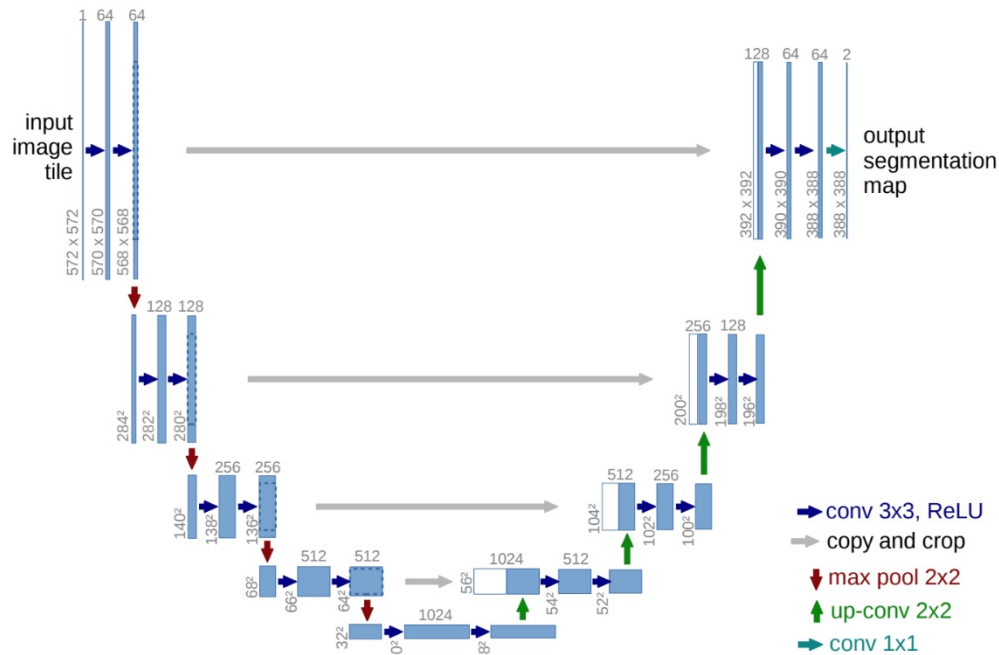
- Dropout regularisation is a simple computational method to prevent over-fitting to data
- SpatialDropout2D in Keras drops (sets to 0) entire feature maps with a given probability during training
- Helps promote independence between feature maps

# Up Convolutions (scaling up)



- Also known as **transpose convolutions**
- Max pooling scales an image down ( $\frac{1}{2} \times \frac{1}{2}$  smaller)
- Up Convolutions scale an image up (e.g.  $2 \times 2$  bigger)
- Want 4 pixels to become 16
- Roughly (??) it works by:
  - Rearrange  $2 \times 2$  input to be  $4 \times 1$
  - Rearrange a  $3 \times 3$  kernel into a  $16 \times 4$  by repeating entries
  - The  $16 \times 4$  then maps a  $4 \times 1$  to a  $1 \times 16$
  - Rearrange  $1 \times 16$  to a  $4 \times 4$  output


# Skip Connections



- Copy some of the outputs from the scaling down to the scaling up side
- Crop middle and copy, add/combine with up side of the network




# UNET Output Segmentation





- Output is a binary image that selects the object region
- May be multiple binary images, one for each object class
- Note that while the input is 572x572 output is 388x388
- Lose some of the edges due to convolutions
- Later versions of U-Net use **padding** for convolutions to maintain original image size

# Padding Convolutions

Suppose we had a 4 x 4 image & we applied a 3x3 convolution



Input 4x4 image




Output 2x2 image

The centre of 3x3 conv can only fit in the **green squares** above


0	0	0	0	0	0
0					0
0					0
0					0
0					0
0	0	0	0	0	0

Padded 6x6 input image



0	0	0	0	0	0
0					0
0					0
0					0
0					0
0	0	0	0	0	0



Output 4x4 image



# U-net



# UNET Advantages and Disadvantages

- Fast to segment
- Field-leading accuracy for segmentation
- Do not need large numbers (1000s) of examples
- Does not separate objects of same type
- Labor intensive to create manual segmentations for training
  - Every training image needs to be completely segmented
  - Conversely, a basic pixel classifier approach does not need this
- Now implemented in Keras / tensorflow for easy use

# Real-World Applications of U-Net



## Medical images

<https://arxiv.org/pdf/2011.01118.pdf>



## Ore segmentation

<https://pubs.rsc.org/en/content/articlelanding/2020/ra/c9ra05877j>



## Satellite Images

<https://arxiv.org/pdf/2003.02899.pdf>

# U-Net in Keras

- Three U-net networks are built into Keras
  - Vanilla: based in the original implementation of U-Net
  - Custom: a customisable U-Net architecture**
  - Satellite: optimised for satellite imaging
- There are also several utility functions to help with training and visualisation of the data and model outputs
- Documentation: <https://pypi.org/project/keras-unet/>
- Keras U-Net model python code:  
[https://github.com/karolzak/keras-unet/tree/master/keras\\_unet/models](https://github.com/karolzak/keras-unet/tree/master/keras_unet/models)



# Some UNET Resources

- A line by line explanation/construction of UNET in tensorflow  
<https://towardsdatascience.com/unet-line-by-line-explanation-9b191c76baf5>
- The original U-net paper  
<https://arxiv.org/pdf/1505.04597.pdf>
- U-net in Keras  
<https://pypi.org/project/keras-unet/>
- Keras documentation on layer types  
Convolution: [https://keras.io/api/layers/convolution\\_layers/convolution2d/](https://keras.io/api/layers/convolution_layers/convolution2d/)  
Max Pooling: [https://keras.io/api/layers/pooling\\_layers/max\\_pooling2d/](https://keras.io/api/layers/pooling_layers/max_pooling2d/)  
Conv Transpose: [https://keras.io/api/layers/convolution\\_layers/convolution2d\\_transpose/](https://keras.io/api/layers/convolution_layers/convolution2d_transpose/)  
Dropout: [https://keras.io/api/layers/regularization\\_layers/dropout/](https://keras.io/api/layers/regularization_layers/dropout/)
- An explanation of Up Sampling  
<https://naokishibuya.medium.com/up-sampling-with-transposed-convolution-9ae4f2df52d0>

# Data Preparation, Training UNET and Evaluating the Results

# Creating a Training Set

- In to create a U-net segmenter, you will need to create a set of labeled examples
- This represents the **ground truth** that U-net will learn
- This is typically a laborious and boring process drawing outlines on numerous images



- How many you will need will depend on the problem
- **Data augmentation** during training can reduce the number

# Data Augmentation for Imaging

- Creating image sets is time consuming and there may be limited numbers of examples available
- One approach to extending a set of images is to apply geometric transformations to create new examples




Original



Flip horizontal



Flip vertical



Rotate



Shear

- Keras has **ImageDataGenerator** and **flow\_from\_directory**
- Keras-Unet has **get\_augmented** function

# Image Batch Generators for Training

- Usually, image sets are too large to present all at once during training as they will not fit in CPU/GPU memory
- Also calculations to update network weights may get too large
- Hence training usually occurs in **batches**, where subsets of images of a fixed size are selected and trained on
- Subsets are repeatedly drawn from the training set, and the network trained, until all of the training set had been presented
- The process is usually repeated multiple times
- In Keras you can set **batchsize** as a parameter to **model.fit** and pass all of your training data to **model.fit** (memory expensive)  
Or you can use one of the batch generators, and pass that to the **model.fit** function (we will see examples in the code).

# Using Google Colab

- Colab is a web-based iPython Notebook service
- You can play/run blocks of python code interactively
- Your code can run on Googles CPU/GPU or TPU in the cloud
- It is free but has usage limitations
- Using Tensorflow with image applications on Colab with GPU can be 10 or 20 times faster than with CPU  
( In Colab set Runtime / Change Runtime Type → GPU)
- TPUs are Tensorflow Processing units – designed for Tensorflow
- Colab Code is interchangeable with Jupyter Notebooks
- Colab basics:

[https://colab.research.google.com/drive/1YKHHLSIG-B9Ez2-zf-YFxXTVgfC\\_Aqtt](https://colab.research.google.com/drive/1YKHHLSIG-B9Ez2-zf-YFxXTVgfC_Aqtt)



# Setting up Google Colab for the Workshop

- Create a folder in your Google Drive called **Colab Notebooks**
- Download code and slides for this workshop:  
<https://github.com/doktor-nick/intro-to-ml-for-imaging/archive/master.zip>
- Unzip it and copy the folder into **Colab Notebooks**
- Go into the folder **intro-to-ml-for-imaging-master**
- Right-click on [Unet\\_in\\_Keras\\_introduction.ipynb](#)
- Select **Open with / + Connect more apps**
- Search for and install Colaboratory
- Double clicking on the [Unet\\_in\\_Keras\\_introduction.ipynb](#) should open it in Colab
- When running the code for the workshop you will need to link and authorise your Google Drive to be used by the notebook. The provided notebook has code/instructions for doing this.

# **Hands On**

Unet\_in\_Keras\_introduction.ipynb  
Colab Notebook

Optimisers, loss functions, and tiling

# Measuring Success: Loss functions for imaging

- A common measure of error/loss in machine learning is root mean square error (RMSE), i.e. the sum of the square of the differences between the real and predicted values
- But predicting a segmentation is a **binary** problem,  
1 = in object, 0 = not in object
- **Binary Cross-entropy** a bit is like RMSE, but it is more punitive on getting a prediction in the wrong class.
- The formula looks like this


$$H_p(q) = -\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i))$$

More explanation here:


<https://towardsdatascience.com/understanding-binary-cross-entropy-log-loss-a-visual-explanation-a3ac6025181a>

# Unbalanced problems and loss functions

- One problem with RMSE and binary cross-entropy is that they are not good when the problem is **unbalanced**, that is if there is a lot more of one class than another.
- For instance,



Target




Prediction


the RMSE between these two images is actually quite small because they agree a lot on the background

# Jaccard loss : “intersection over union”

- One problem with RMSE and binary cross-entropy is that they are not good when the problem is **unbalanced**, that is if there is a lot more of one class than another.
- Jaccard: intersection over union  $(|X \cap Y|) / (|X| + |Y| - |X \cap Y|)$



object prediction




Agree  
-----  
Disagree + Agree

- More discussion on image metrics:  
<https://towardsdatascience.com/image-segmentation-choosing-the-correct-metric-aa21fd5751af>

# Image Tiling

- Available memory on a GPU can often be an issue for imaging
- During training, **reducing batch** size may help or **scaling down**
- **Tiling**, that is chopping an image into smaller pieces is another approach
- Care needs to be taken that each tile contains enough information so that it can be accurately segmented



- Usually want to use **overlapping tiles** to avoid edge effects
- The Keras Unet libraries can generate patches easily

# Which Optimiser for Semantic Image Segmentation?

- In Keras there are a range of optimisers available including SGD, RMSprop, Adam, Adadelata, Adagrad, Adamax, Nadam Ftrl
- Performance may vary according to the application
- Generally Adam and SGD are reasonable choices
- For a discussion of different optimiser types see <https://runder.io/optimizing-gradient-descent/>
- For optimisers available in Keras see <https://keras.io/api/optimizers/>



# Overview of choices in network/evaluations

- Optimiser: Adam
- Loss function: Jaccard distance
- Reduce batchsize, scale images down, or tile if size is a problem
- Sometimes the optimiser will get stuck. It is worth running the fitting functions more than once to ensure better results

**Hands on**

Unet\_in\_Keras\_further\_topics.ipynb  
Colab Notebook

**The End**