



Datasets and slides:

<https://github.com/doktor-nick/adv-ML-image-segmentation-with-UNET>

Advanced Machine Learning Image Segmentation with UNET



Dr Nick Hamilton
Institute Bio-Mathematician
Institute for Molecular Bioscience
Research Computing Centre
Queensland Cyber Infrastructure Foundation
n.hamilton@imb.uq.edu.au



www.imb.uq.edu.au

Power corrupts.

PowerPoint corrupts absolutely.

Outline

- Brief machine vision applications overview
- Semantic image segmentation
- The UNET architecture for image segmentation
- Setting up Google Colab
- Exercises
 - Introduction for UNET
 - More advanced UNET use


Applications of Computer Vision

- Classification



Rainbow bee-eater

- Object Detection



Bee-eater, bee

- Semantic Segmentation




Bee-eater, bee, stick,
background

- Instance Segmentation



Bee-eater 1 and 2

Image Classification



=

```
( 0.000772532, 0.00154506, 0.00343348, 0.0139056, 0.0308155,  
 0.0440343, 0.0337339, 0.0433476, 0.828412, 0.000235031,  
 0.000352547, 0.00205653, 0.00593454, 0.0223867, 0.0351372,  
 0.022093, 0.0312004, 0.880604, 0.000420663, 0.00117786,  
 0.00496382, 0.0165741, 0.058388, 0.107017, 0.0904425,  
 0.110803, 0.61024)
```



Cell
Cytoskeleton

Convolutional networks are typically good at extracting features for classification.

See “Introduction for Deep Learning and Tensorflow” Workshop.

Image Segmentation



Semantic Image Segmentation

Definition

A segmentation is a partition of an image into regions that label the class of each region. e.g. foreground/background

Approach

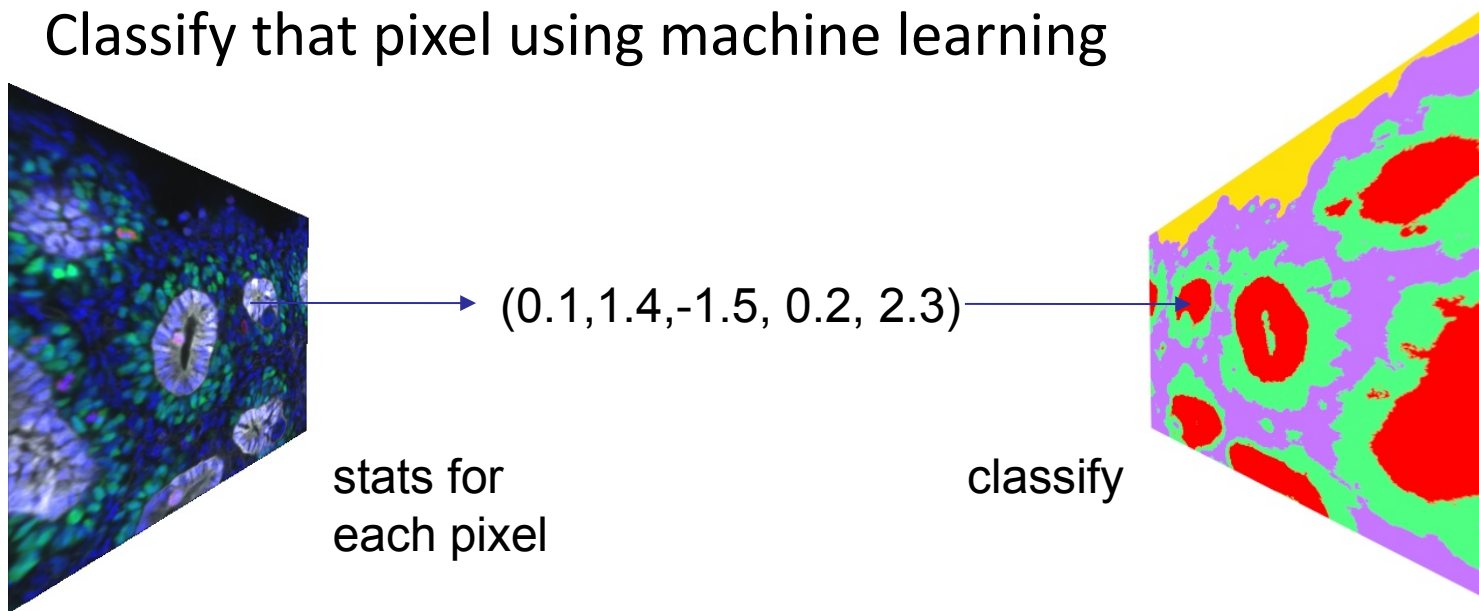
Instead of a whole image we want to label/classify individual pixels



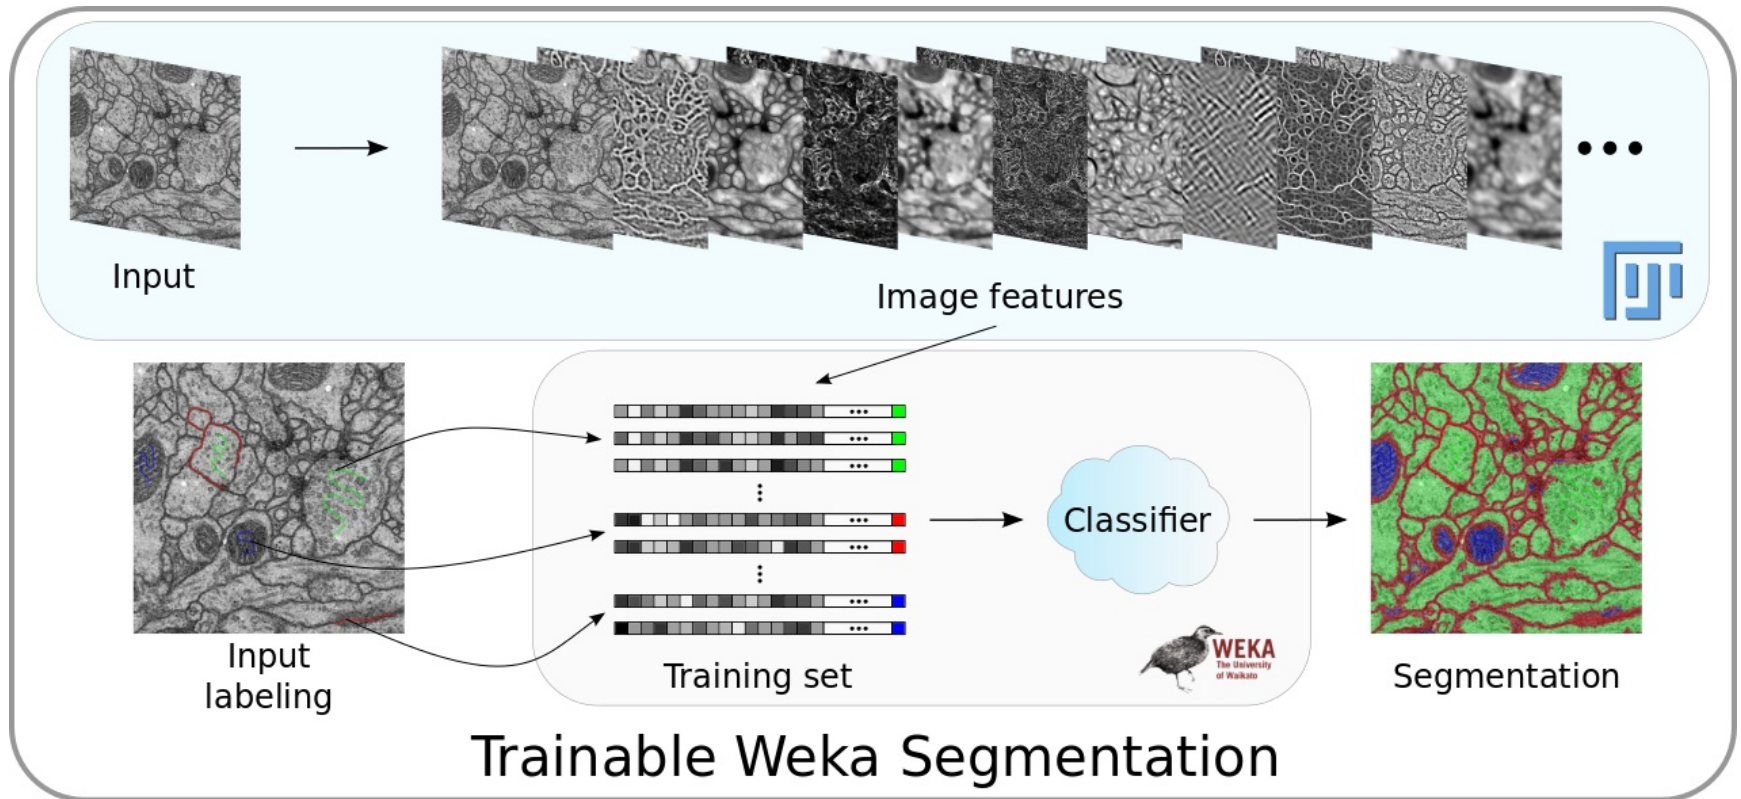
Ureteric tree
Cap mesenchyme
Stroma
Background

Basic Image Segmentation

- Generate **statistics** for **each pixel** in the image
- E.g. intensity, local median intensity, gradient, ...
- Classify that pixel using machine learning



The ImageJ/Fiji Trainable Weka Pipeline




See “Introduction to Machine Learning for Imaging” workshop

Image: https://imagej.net/Trainable_Weka_Segmentation

The UNET Architecture for Image Segmentation

Convolutional Deep Learning with U-net

- **U-net** was a new type of neural network architecture introduced in 2015
- It down-scales then up-scales an image to create a segmentation




- U-Net and its variants are the state of the art in segmentation
- U-net paper: <https://arxiv.org/pdf/1505.04597.pdf>

Principle UNET Components

- 3x3 **Convolution** followed by ReLU (rectified linear activation function)
- 2 x 2 **Maxpooling** with stride 2 (halves image size)
- Double number of feature channels (convolutions) 64, 128, 256, 512, at each down sampling to compensate
- **Dropout layers**
- **Skip layers**
- **Transpose convolutions** / up convolutions

Convolutions



$$Output = \sum Kernel_{i,j} * Input_{i,j}$$

- The kernel represents the pattern to be detected
- The more the input matches the kernel, the more positive the output response would be
- Convolutions are pattern detectors
- The kernel weights are learned during training
i.e. it learns what patterns to detect and respond to

A 3x3 Convolution Max

- **Convolution** of an image by a 3x3 matrix replaces each pixel in an image by a weighted sum of those pixels adjacent to it, i.e.

$$\begin{aligned} I'(x,y) = & a I(x-1,y+1) + b I(x,y+1) + c I(x+1,y+1) \\ & d I(x-1,y) + e I(x,y) + f I(x+1,y) \\ & g I(x-1,y-1) + h I(x,y-1) + i I(x+1,y-1) \end{aligned}$$

Local Image Region

$I(x-1,y+1)$	$I(x,y+1)$	$I(x+1,y+1)$
$I(x-1,y)$	$I(x,y)$	$I(x+1,y)$
$I(x-1,y-1)$	$I(x,y-1)$	$I(x+1,y-1)$

Convolution matrix

a	b	c
d	e	f
g	h	i

Convolution Example

Convolution

0	0	0
0	1	1
0	1	1

Image

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	1	1	0	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

Image x Convolution

Dark	Light	Dark	Black	Black
Light	White	Light	Black	Black
Dark	Light	Dark	Black	Black
Black	Black	Black	Black	Black
Black	Black	Black	Black	Black

Max Pooling with step size 2

10	6	6	0
5	-3	14	-2
0	2	-7	-3
0	0	-5	-5

Input 4x4 image



10	14
2	-3


Output 2x2 image

- Slide a window across the input and pick a value at every window position.
- Max pooling – take the max value.
- Average pooling – take the average value.
- Pooling layers are information filters.
- Images are reduced in scale, typically by $\frac{1}{2} \times \frac{1}{2}$

Spatial Dropout Layers

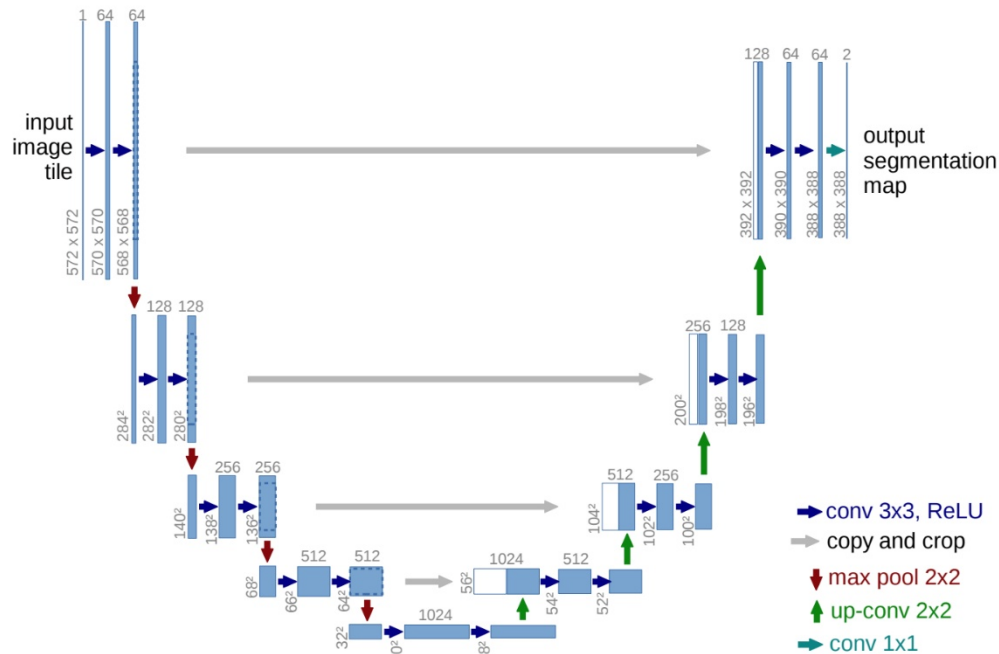
- Dropout regularisation is a simple computational method to prevent over-fitting to data
- SpatialDropout2D in Keras drops (sets to 0) entire feature maps with a given probability during training
- Helps promote independence between feature maps

Up Convolutions (scaling up)




- Also known as **transpose convolutions**
- Max pooling scales an image down ($\frac{1}{2} \times \frac{1}{2}$ smaller)
- Up Convolutions scale an image up (e.g. 2×2 bigger)
- Want 4 pixels to become 16
- Roughly (??) it works by:
 - Rearrange 2×2 input to be 4×1
 - Rearrange a 3×3 kernel into a 16×4 by repeating entries
 - The 16×4 then maps a 4×1 to a 1×16
 - Rearrange 1×16 to a 4×4 output

Skip Connections



- Copy some of the outputs from the scaling down to the scaling up side
- Crop middle and copy, add/combine with up side of the network


UNET Output Segmentation





- Output is a binary image that selects the object region
- May be multiple binary images, one for each object class
- Note that while the input is 572x572 output is 388x388
- Lose some of the edges due to convolutions
- Later versions of U-Net use **padding** for convolutions to maintain original image size

Padding Convolutions

Suppose we had a 4 x 4 image & we applied a 3x3 convolution



Input 4x4 image




Output 2x2 image

The centre of 3x3 conv can only fit in the **green squares** above


0	0	0	0	0	0
0					0
0					0
0					0
0					0
0	0	0	0	0	0

Padded 6x6 input image



0	0	0	0	0	0
0					0
0					0
0					0
0					0
0	0	0	0	0	0



Output 4x4 image



U-net



UNET Advantages and Disadvantages

- Fast to segment
- Field-leading accuracy for segmentation
- Do not need large numbers (1000s) of examples
- Does not separate objects of same type
- Labor intensive to create manual segmentations for training
 - Every training image needs to be completely segmented
 - Conversely, a basic pixel classifier approach does not this
- Now implemented in Keras / tensorflow for easy use

Real-World Applications of U-Net



Medical images

<https://arxiv.org/pdf/2011.01118.pdf>



Ore segmentation

<https://pubs.rsc.org/en/content/articlelanding/2020/ra/c9ra05877j>



Satellite Images

<https://arxiv.org/pdf/2003.02899.pdf>

U-Net in Keras

- Three U-net networks are built into Keras
 - Vanilla: based in the original implementation of U-Net
 - Custom: a customisable U-Net architecture**
 - Satellite: optimised for satellite imaging
- There are also several utility functions to help with training and visualisation of the data and model outputs
- Documentation: <https://pypi.org/project/keras-unet/>
- Keras U-Net model python code:
https://github.com/karolzak/keras-unet/tree/master/keras_unet/models

Some UNET Resources

- A line by line explanation/construction of UNET in tensorflow
<https://towardsdatascience.com/unet-line-by-line-explanation-9b191c76baf5>
- The original U-net paper
<https://arxiv.org/pdf/1505.04597.pdf>
- U-net in Keras
<https://pypi.org/project/keras-unet/>
- Keras documentation on layer types
Convolution: https://keras.io/api/layers/convolution_layers/convolution2d/
Max Pooling: https://keras.io/api/layers/pooling_layers/max_pooling2d/
Conv Transpose: https://keras.io/api/layers/convolution_layers/convolution2d_transpose/
Dropout: https://keras.io/api/layers/regularization_layers/dropout/
- An explanation of Up Sampling
<https://naokishibuya.medium.com/up-sampling-with-transposed-convolution-9ae4f2df52d0>

Data Preparation, Training UNET and Evaluating the Results

Creating a Training Set

- In to create a U-net segmenter, you will need to create a set of labeled examples
- This represents the **ground truth** that U-net will learn
- This is typically a laborious and boring process drawing outlines on numerous images



- How many you will need will depend on the problem
- **Data augmentation** during training can reduce the number

Data Augmentation for Imaging

- Creating image sets is time consuming and there may be limited numbers of examples available
- One approach to extending a set of images is to apply geometric transformations to create new examples




Original




Flip horizontal



Flip vertical



Rotate



Shear

- Keras has **ImageDataGenerator** and **flow_from_directory**
- Keras-Unet has **get_augmented** function

Image Batch Generators for Training

- Usually, image sets are too large to present all at once during training as they will not fit in CPU/GPU memory
- Also calculations to update network weights may get too large
- Hence training usually occurs in **batches**, where subsets of images of a fixed size are selected and trained on
- Subsets are repeatedly drawn from the training set, and the network trained, until all of the training set had been presented
- The process is usually repeated multiple times
- In Keras you can set **batchsize** as a parameter to **model.fit** and pass all of your training data to **model.fit** (memory expensive)
Or you can use one of the batch generators, and pass that to the **model.fit** function (we will see examples in the code).

Using Google Colab

- Colab is a web-based iPython Notebook service
- You can play/run blocks of python code interactively
- Your code can run on Googles CPU/GPU or TPU in the cloud
- It is free but has usage limitations
- Using Tensorflow with image applications on Colab with GPU can be 10 or 20 times faster than with CPU
(In Colab set Runtime / Change Runtime Type → GPU)
- TPUs are Tensorflow Processing units – designed for Tensorflow
- Colab Code is interchangeable with Jupyter Notebooks
- Colab basics:

https://colab.research.google.com/drive/1YKHHLSIG-B9Ez2-zf-YFxXTVgfC_Aqtt

Setting up Google Colab for the Workshop

- Create a folder in your Google Drive called **Colab Notebooks**
- Download code and slides for this workshop:
<https://github.com/doktor-nick/intro-to-ml-for-imaging/archive/master.zip>
- Unzip it and copy the folder into **Colab Notebooks**
- Go into the folder **intro-to-ml-for-imaging-master**
- Right-click on [Unet_in_Keras_introduction.ipynb](#)
- Select **Open with / + Connect more apps**
- Search for and install Colaboratory
- Double clicking on the [Unet_in_Keras_introduction.ipynb](#) should open it in Colab
- When running the code for the workshop you will need to link and authorise your Google Drive to be used by the notebook. The provided notebook has code/instructions for doing this.

Hands On

Unet_in_Keras_introduction.ipynb
Colab Notebook

Optimisers, loss functions, and tiling

Measuring Success: Loss functions for imaging

- A common measure of error/loss in machine learning is root mean square error (RMSE), i.e. the sum of the square of the differences between the real and predicted values
- But predicting a segmentation is a **binary** problem,
1 = in object, 0 = not in object
- Binary Cross-entropy a bit is like RMSE, but it is more punitive on getting a prediction in the wrong class.
- The formula looks like this


$$H_p(q) = -\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i))$$

More explanation here:


<https://towardsdatascience.com/understanding-binary-cross-entropy-log-loss-a-visual-explanation-a3ac6025181a>

Unbalanced problems and loss functions

- One problem with RMSE and binary cross-entropy is that they are not good when the problem is **unbalanced**, that is if there is a lot more of one class than another.
- For instance,



Target




Prediction


the RMSE between these two images is actually quite small because they agree a lot on the background

Jaccard loss : “intersection over union”

- One problem with RMSE and binary cross-entropy is that they are not good when the problem is **unbalanced**, that is if there is a lot more of one class than another.
- Jaccard: intersection over union $(|X \cap Y|) / (|X| + |Y| - |X \cap Y|)$



object prediction




$$\frac{\text{Agree}}{\text{Agree} + \text{Agree}}$$

- More discussion on image metrics:
<https://towardsdatascience.com/image-segmentation-choosing-the-correct-metric-aa21fd5751af>

Image Tiling

- Available memory on a GPU can often be an issue for imaging
- During training, **reducing batch** size may help or **scaling down**
- **Tiling**, that is chopping an image into smaller pieces is another approach
- Care needs to be taken that each tile contains enough information so that it can be accurately segmented



- Usually want to use **overlapping tiles** to avoid edge effects
- The Keras Unet libraries can generate patches easily

Which Optimiser for Semantic Image Segmentation?

- In Keras there are a range of optimisers available including SGD, RMSprop, Adam, Adadelata, Adagrad, Adamax, Nadam Ftrl
- Performance may vary according to the application
- Generally Adam and SGD are reasonable choices
- For a discussion of different optimiser types see <https://runder.io/optimizing-gradient-descent/>
- For optimisers available in Keras see <https://keras.io/api/optimizers/>

Overview of choices in network/evaluations

- Optimiser: Adam
- Loss function: Jaccard distance
- Reduce batchsize, scale images down, or tile if size is a problem
- Sometimes the optimiser will get stuck. It is worth running the fitting functions more than once to ensure better results

Hands on

Unet_in_Keras_further_topics.ipynb
Colab Notebook

The End