



Datasets and slides:

<https://github.com/doktor-nick/adv-ML-image-segmentation-with-UNET>

# Advanced Machine Learning Image Segmentation with UNET



**Dr Nick Hamilton**  
**Institute Bio-Mathematician**  
Institute for Molecular Bioscience  
Research Computing Centre  
Queensland Cyber Infrastructure Foundation  
[n.hamilton@imb.uq.edu.au](mailto:n.hamilton@imb.uq.edu.au)



[www.imb.uq.edu.au](http://www.imb.uq.edu.au)

Power corrupts.

PowerPoint corrupts absolutely.

# Outline

- Brief machine vision applications overview
- Semantic image segmentation
- The UNET architecture for image segmentation
- Setting up Google Colab
- Exercises
  - Introduction for UNET
  - More advanced UNET use


# Applications of Computer Vision

- Classification



Rainbow bee-eater

- Object Detection



Bee-eater, bee

- Semantic Segmentation




Bee-eater, bee, stick,  
background

- Instance Segmentation



Bee-eater, bee

# Image Classification



=

```
( 0.000772532, 0.00154506, 0.00343348, 0.0139056, 0.0308155,  
 0.0440343, 0.0337339, 0.0433476, 0.828412, 0.000235031,  
 0.000352547, 0.00205653, 0.00593454, 0.0223867, 0.0351372,  
 0.022093, 0.0312004, 0.880604, 0.000420663, 0.00117786,  
 0.00496382, 0.0165741, 0.058388, 0.107017, 0.0904425,  
 0.110803, 0.61024)
```



Cell  
Cytoskeleton

Convolutional networks are typically good at extracting features for classification.

See “Introduction for Deep Learning and Tensorflow” Workshop.

# Image Segmentation




# Semantic Image Segmentation

## Definition

A segmentation is a partition of an image into regions that label the class of each region. e.g. foreground/background

## Approach

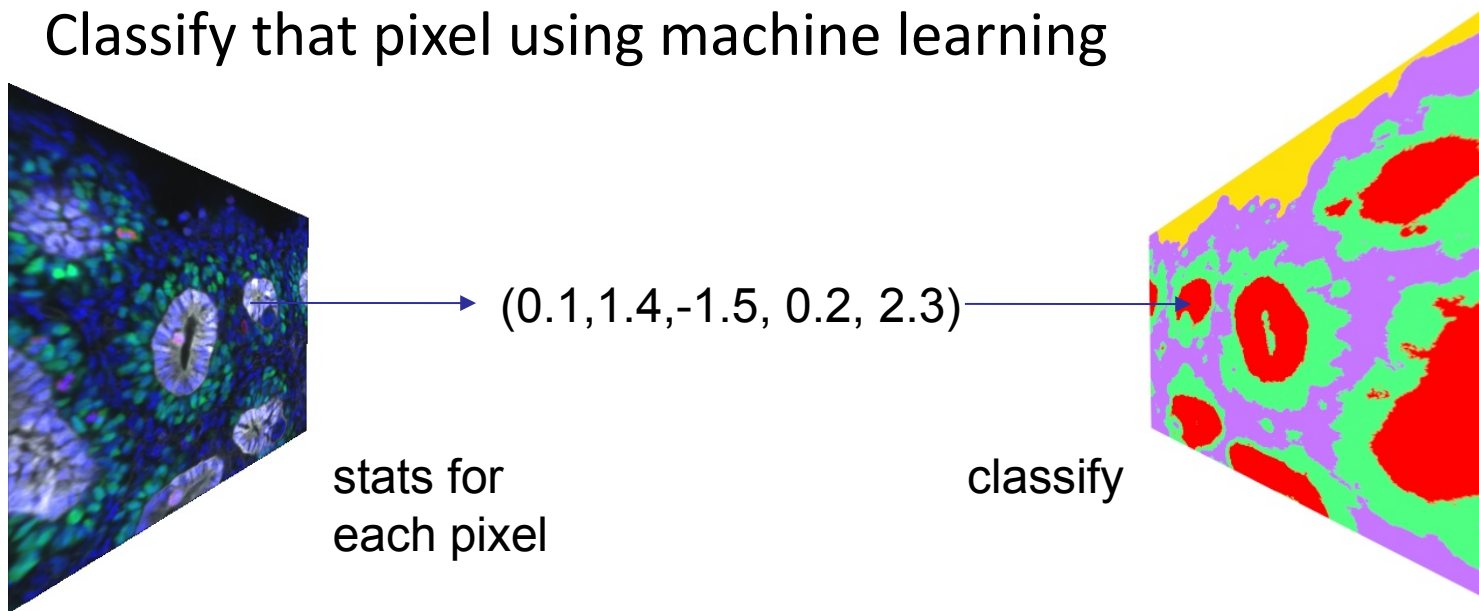
Instead of a whole image we want to label/classify individual pixels




Ureteric tree  
Cap mesenchyme  
Stroma  
Background

# Basic Image Segmentation

- Generate **statistics** for **each pixel** in the image
- E.g. intensity, local median intensity, gradient, ...
- Classify that pixel using machine learning



# The ImageJ/Fiji Trainable Weka Pipeline



See “Introduction to Machine Learning for Imaging” workshop


Image: [https://imagej.net/Trainable\\_Weka\\_Segmentation](https://imagej.net/Trainable_Weka_Segmentation)



# The UNET Architecture for Image Segmentation

# Convolutional Deep Learning with U-net

- **U-net** was a new type of neural network architecture introduced in 2015
- It down-scales then up-scales an image to create a segmentation




- U-Net and its variants are the state of the art in segmentation
- U-net paper: <https://arxiv.org/pdf/1505.04597.pdf>

# Principle UNET Components

- 3x3 **Convolution** followed by ReLU (rectified linear activation function)
- 2 x 2 **Maxpooling** with stride 2 (halves image size)
- Double number of feature channels (convolutions) 64, 128, 256, 512, at each down sampling to compensate
- **Dropout layers**
- **Skip layers**
- **Transpose convolutions** / up convolutions

# Convolutions



$$Output = \sum Kernel_{i,j} * Input_{i,j}$$

- The kernel represents the pattern to be detected
- The more the input matches the kernel, the more positive the output response would be
- Convolutions are pattern detectors
- The kernel weights are learned during training  
i.e. it learns what patterns to detect and respond to

# Max Pooling with step size 2

10	6	6	0
5	-3	14	-2
0	2	-7	-3
0	0	-5	-5

Input 4x4 image



10	14
2	-3

Output 2x2 image


- Slide a window across the input and pick a value at every window position.
- Max pooling – take the max value.
- Average pooling – take the average value.
- Pooling layers are information filters.
- Images are reduced in scale, typically by  $\frac{1}{2} \times \frac{1}{2}$



# Spatial Dropout Layers

- Dropout regularisation is a simple computational method to prevent over-fitting to data
- SpatialDropout2D in Keras drops (sets to 0) entire feature maps with a given probability during training
- Helps promote independence between feature maps

# Up Convolutions (scaling up)



- Also known as **transpose convolutions**
- Max pooling scales an image down ( $\frac{1}{2} \times \frac{1}{2}$  smaller)
- Up Convolutions scale an image up (e.g.  $2 \times 2$  bigger)
- Want 4 pixels to become 16
- Roughly (??)

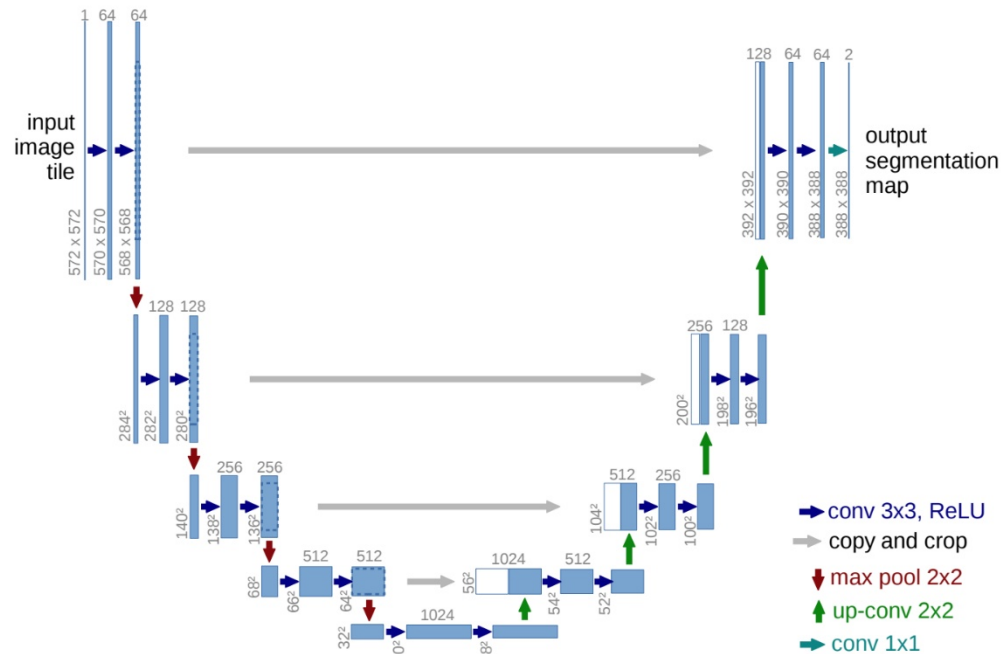
Rearrange 2x2 input to be 4x1

Rearrange a 3x3 kernel into a 16x4 by repeating entries

The 16x4 then maps a 4x1 to a 1x16


Rearrange 1x16 to a 4x4 output

# Skip Connections



- Copy some of the outputs from the scaling down to the scaling up side
- Crop middle and copy, add/combine with up side of the network


# UNET Output Segmentation




- Output is a binary image that selects the object region
- May be multiple binary images, one for each object class
- Note that while the input is 572x572 output is 388x388
- Lose some of the edges due to convolutions
- Later versions of U-Net use **padding** for convolutions to maintain original image size

# Padding Convolutions

Suppose we had a 4 x 4 image & we applied a 3x3 convolution



Input 4x4 image




Output 2x2 image

The centre of 3x3 conv can only fit in the **green squares** above

0	0	0	0	0	0
0					0
0					0
0					0
0					0
0	0	0	0	0	0

Padded 6x6 input image

0	0	0	0	0	0
0					0
0					0
0					0
0					0
0	0	0	0	0	0




Output 4x4 image





# UNET Advantages and Disadvantages

- Fast to segment
- Field-leading accuracy for segmentation
- Do not need large numbers (1000s) of examples
- Does not separate objects of same type
- Labor intensive to create manual segmentations for training
  - Every training image needs to be completely segmented
  - Conversely, a basic pixel classifier approach does not this
- Now implemented in Keras / tensorflow for easy use




# U-net



# Real-World Applications of U-Net



## Medical images

<https://arxiv.org/pdf/2011.01118.pdf>



## Ore segmentation

<https://pubs.rsc.org/en/content/articlelanding/2020/ra/c9ra05877j>



## Satellite Images

<https://arxiv.org/pdf/2003.02899.pdf>

# U-Net in Keras

- Three U-net networks are built into Keras
  - Vanilla: based in the original implementation of U-Net
  - Custom: a customisable U-Net architecture
  - Satellite: optimised for satellite imaging
- There are also several utility functions to help with training and visualisation of the data and model outputs
- Documentation: <https://pypi.org/project/keras-unet/>
- Keras U-Net model python code:  
[https://github.com/karolzak/keras-unet/tree/master/keras\\_unet/models](https://github.com/karolzak/keras-unet/tree/master/keras_unet/models)

# Some UNET Resources

- A line by line explanation/construction of UNET in tensorflow  
<https://towardsdatascience.com/unet-line-by-line-explanation-9b191c76baf5>
- The original U-net paper  
<https://arxiv.org/pdf/1505.04597.pdf>
- U-net in Keras  
<https://pypi.org/project/keras-unet/>
- Keras documentation on layer types  
Convolution: [https://keras.io/api/layers/convolution\\_layers/convolution2d/](https://keras.io/api/layers/convolution_layers/convolution2d/)  
Max Pooling: [https://keras.io/api/layers/pooling\\_layers/max\\_pooling2d/](https://keras.io/api/layers/pooling_layers/max_pooling2d/)  
Conv Transpose: [https://keras.io/api/layers/convolution\\_layers/convolution2d\\_transpose/](https://keras.io/api/layers/convolution_layers/convolution2d_transpose/)  
Dropout: [https://keras.io/api/layers/regularization\\_layers/dropout/](https://keras.io/api/layers/regularization_layers/dropout/)
- An explanation of Up Sampling  
<https://naokishibuya.medium.com/up-sampling-with-transposed-convolution-9ae4f2df52d0>



# Data Preparation, Training UNET and Evaluating the Results

# Creating a Training Set

- In to create a U-net segmenter, you will need to create a set of labeled examples
- This represents the **ground truth** that U-net will learn
- This is typically a laborious and boring process drawing outlines on numerous images



- How many you will need will depend on the problem
- **Data augmentation** during training can reduce the number

# Data Augmentation for Imaging

- Creating image sets is time consuming and there may be limited numbers of examples available
- One approach to extending a set of images is to apply geometric transformations to create new examples




Original




Flip horizontal



Flip vertical



Rotate



Shear

- Keras has **ImageDataGenerator** and **flow\_from\_directory**
- Keras-Unet has **get\_augmented** function

# Image Batch Generators for Training

- Usually, image sets are too large to present all at once during training as they will not fit in CPU/GPU memory
- Also calculations to update network weights may get too large
- Hence training usually occurs in **batches**, where subsets of images of a fixed size are selected and trained on
- Subsets are repeatedly drawn from the training set, and the network trained, until all of the training set had been presented
- The process is usually repeated multiple times
- In Keras you can set **batchsize** as a parameter to **model.fit** and pass all of your training data to **model.fit**
- The disadvantage of this is that it makes data **augmentation** tricky ...

# Using Google Colab

- Colab is a web-based iPython Notebook service
- You can play/run blocks of python code interactively
- Your code can run on Googles CPU/GPU or TPU in the cloud
- It is free but has usage limitations
- Using Tensorflow with image applications on Colab with GPU can be 10 or 20 times faster than with CPU  
( In Colab set Runtime / Change Runtime Type → GPU)
- TPUs are Tensorflow Processing units – designed for Tensorflow
- Colab Code is interchangeable with Jupyter Notebooks
- Colab basics:

[https://colab.research.google.com/drive/1YKHHLSIG-B9Ez2-zf-YFxXTVgfC\\_Aqtt](https://colab.research.google.com/drive/1YKHHLSIG-B9Ez2-zf-YFxXTVgfC_Aqtt)



# Setting up Google Colab for the Workshop

- Create a folder in your Google Drive called **Colab Notebooks**
- Download code and slides for this workshop:  
<https://github.com/doktor-nick/intro-to-ml-for-imaging/archive/master.zip>
- Unzip it and copy the folder into **Colab Notebooks**
- Go into the folder **intro-to-ml-for-imaging-master**
- Right-click on [Unet\\_in\\_Keras\\_introduction.ipynb](#)
- Select **Open with / + Connect more apps**
- Search for and install Colaboratory
- Double clicking on the [Unet\\_in\\_Keras\\_introduction.ipynb](#) should open it in Colab
- When running the code for the workshop you will need to link and authorise your Google Drive to be used by the notebook. The provided notebook has code/instructions for doing this.

# **Hands On**

Unet\_in\_Keras\_introduction.ipynb  
Colab Notebook


# Measuring Success: Loss functions for imaging

- Binary Cross-entropy
- Jaccard : intersection over union.  
$$(|X \cap Y|) / (|X| + |Y| - |X \cap Y|)$$

Good for unbalanced problems.
- Other metrics, cf `metrics.py`

# Image Tiling

- Available memory on a GPU can often be an issue for imaging
- During training, **reducing batch** size may help or **scaling down**
- **Tiling**, that is chopping an image into smaller pieces is another approach
- Care needs to be taken that each tile contains enough information so that it can be accurately segmented



- Usually want to use **overlapping tiles** to avoid edge effects
- The Keras Unet libraries can generate patches easily

# Which Optimiser for Semantic Image Segmentation?

- In Keras there are a range of optimisers available including SGD, RMSprop, Adam, Adadelata, Adagrad, Adamax, Nadam Ftrl
- Performance may vary according to the application
- Generally Adam and SGD are reasonable choices
- For a discussion of different optimiser types see <https://runder.io/optimizing-gradient-descent/>
- For optimisers available in Keras see <https://keras.io/api/optimizers/>

# Overview of choices in network/evaluations

- 
- Optimiser: Adam
- Loss function: Jaccard distance
- Batchsize if size is a problem
- 
- Sometimes the optimiser will get stuck. It is worth running the fitting functions more than once to ensure better results
-

**Hands on**

Unet\_in\_Keras\_further\_topics.ipynb  
Colab Notebook

# Random notes 1

- Colab getting started  
<https://towardsdatascience.com/getting-started-with-google-colab-f2fff97f594c>
- Colab overview
- [https://colab.research.google.com/drive/1CRYG\\_y2ACXy0Q3NJcsbNHhF4658jvVox](https://colab.research.google.com/drive/1CRYG_y2ACXy0Q3NJcsbNHhF4658jvVox)
- How to make a colab notebook available to others
- <https://research.google.com/colaboratory/faq.html#:~:text=Colab%20notebooks%20are%20stored%20in,Google%20Drive%20file%20sharing%20instructions>  
.
- Colab test share:  
<https://colab.research.google.com/drive/1G9gtgfhrfkSX95mJfUC5fFil3z7jslsR?usp=sharing>  
- then 'save a copy in drive'
- Do Monash do viz of convolutions?
- The architectures are given in keras\_unet\_master
- Vanilla scales down, custom does not. Padding? padding='valid', padding="same"



# Random notes 2

- Multi-class def custom\_unet(input\_shape, num\_classes=1, ...
- 
- Code will run on colab without GPU , but it is very slow (5 mins per epoch)
- 
- TPU vs GPU? Google claims TPU 15-30x faster